

Priority Queue

Featuring Binary Heap

Overview

- Simple Implementation of priority_queue
- Quick intro to Graph and Tree
- Binary Heap
- priority_queue with Binary Heap

priority_queue

- Queue by value
- Max-in-First-Out

```
int main() {  
    priority_queue<int> pq;  
    pq.push(4);  
    pq.push(20);  
    pq.push(3);  
  
    while (pq.empty() == false) {  
        cout << pq.top() << endl;  
        pq.pop();  
    }  
}
```

20
4
3

V0.1, priority_queue by vector

- Use `vector` to store data
- Push = simply `push_back`
- Top, pop = find the max value and return/erase
- `max_element` return iterator to max element

```
namespace CP {
    template <typename T>
    class priority_queue
    {
        protected:
            std::vector<T> v;
        public:
            bool empty()                { return v.empty(); }
            bool size()                 { return v.size(); }

            void push(const T &e)       { v.push_back(e); }

            T top()                     { return *std::max_element(v.begin(),v.end()); }

            void pop() { v.erase(std::max_element(v.begin(),v.end())); }

    };
}
```

max_element

```
iterator max_element(iterator first, iterator last)
{
    if (first == last) return last;
    iterator largest = first;
    ++first;
    for (; first != last; ++first)
        if (*largest < *first)
            largest = first;
    return largest;
}
```

$O(n)$

V0.1 complexities

push

top()

pop()

V0.2 faster pop, top (and push??)

- v0.1 has many drawback
 - Consecutive call of top is slow (it shouldn't)
 - Both pop and top works almost the same
- v0.2 focus on slower push while keeps pop, top fast
- Make the vector sorted
 - Max will be at the back
 - Fast pop, top

```
namespace CP {  
    template <typename T>  
    class priority_queue {  
    protected:  
        std::vector<T> v;  
    public:  
        bool empty() { return v.empty(); }  
        size_t size() { return v.size(); }  
  
        T& top() { return v[v.size()-1]; }  
  
        void pop() { v.erase(v.end()-1); }  
  
        void push(const T& e) {  
            // do something  
        }  
    };  
}
```

v0.2 push

Complexity

$O(N \log N)$, where $N = \text{std::distance}(\text{first}, \text{last})$ comparisons on average. (until C++11)

$O(N \log N)$, where $N = \text{std::distance}(\text{first}, \text{last})$ comparisons. (since C++11)

- Maintain that the vector is sorted at every push

```
void push(const T& e) {  
    v.push_back(e);  
    std::sort(v.begin(), v.end());  
}
```

$O(n \log n)$

Why Big-Theta,
not Big O?

```
void push(const T& e) {  
    auto it = v.begin();  
    while (it < v.end() && *it <= e)  
        it++;  
    v.insert(it, e);  
}
```

$\Theta(n)$

Why cannot use
Big-Theta?

```
void push(const T& e) {  
    v.insert(std::upper_bound(v.begin(), v.end(), e), e);  
}
```

Upper bound is $O(\log n)$

$O(n)$

Which one is better?


- v0.1 fast push
- v0.2 fast pop, top
- Depends on which operation we use most often
- The real version works like v0.2, we maintain some **rules** of the data that is stored in the vector such that
 - We know where max is (for fast top)
 - Much faster push, a little bit slower pop
 - Using structure call **Binary Heap**

Graph and Tree

Quick introduction

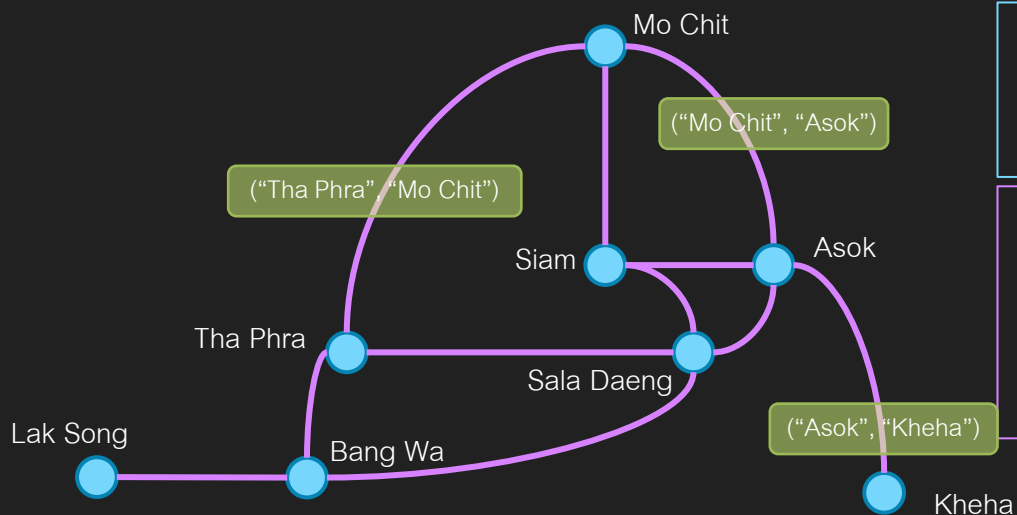
Graph

- Discrete Math Graph
- A math model that describe entities and connectivity between them



Graph Model

- Graph consists of two things
 - Nodes (vertex, vertices) are things we want to connect
 - Edges are pairs, each pair is a connectivity between two node
- Graph $G = (V, E)$ where V is a set of nodes and E is a set of edges




$V = \{ \text{"Mo Chit", "Siam", "Asok", "Sala Daeng", "Tha Phra", "Lak Song", "Bang Wa", "Kheha"} \}$


$E = \{ (\text{"Mo Chit", "Asok"}, (\text{"Mo Chit", "Siam"}, (\text{"Tha Phra", "Mo Chit"}, (\text{"Sala Daeng", "Tha Phra"}), \dots \}$

Tree


- A special kind of graph
 - Has N nodes and $N-1$ Edges
 - Every nodes must be connected (we can start from any node and can walk through edge to reach any node)




Is a tree



Not Tree
(too many edges)



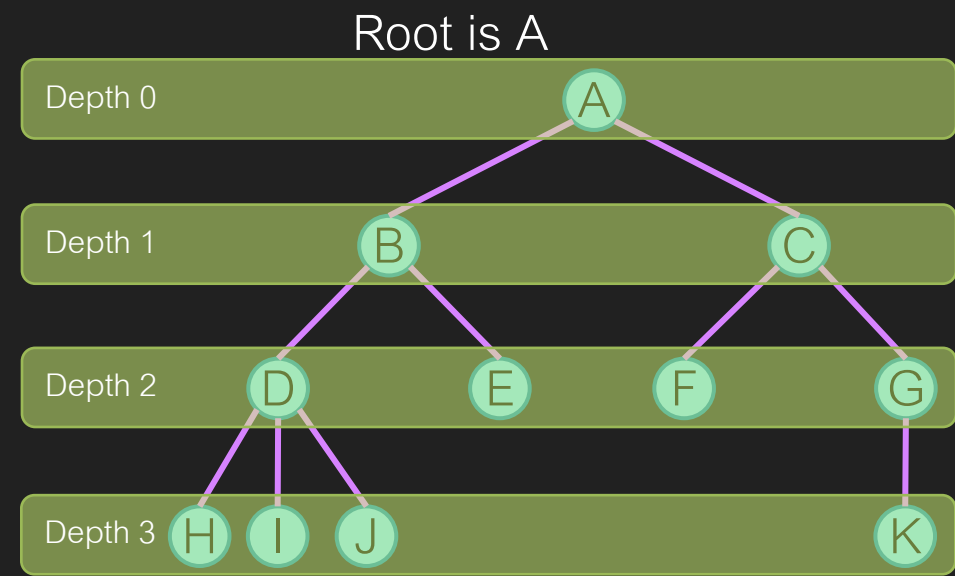
Not Tree
(too many nodes)



Not Tree
(not connected)

Rooted Tree


- Tree where one node is defined as a **root**
- For an edge in a rooted tree, a node that is closer to the root is called **parent** while the other node is called **child**
- **Ancestor** of node A = parent of parent of A
- **Descendant** of node A = child of child of A
- Root is usually drawn at the top and is considered as the **starting point**
 - Root is at level 0 (depth 0)
 - Children of root are drawn at the same level at **level 1** (depth 1)
 - Children of children of root are at **level 2** (depth 2)




A is the parent of B and C
B is the parent of D and E
K is a child of G
H, I and J are children of D
B is an ancestor of D, H, I and J
K is a descendant of G, C, and A

Complete Binary Tree


- Binary Tree = a tree that every node has at most 2 children
- Complete tree = the tree must be filled with every possible node at every level (except the deepest level which must be filled as far to the left as possible)
 - Blank tree is consider a complete binary tree




OK




OK



Not binary



Not complete
(at depth 1)



OK


Exercise

Hint: The answer is unique (There are exactly 1 way to draw a complete binary tree of size k)


- Draw a Complete Binary Tree that has 4, 5, 8, 10 nodes




0 nodes
(blank tree)




1 nodes
(only root)




2 nodes




3 nodes




4 nodes



5 nodes



8 nodes



10 nodes

Special Property of a complete binary tree


- There is exactly one way to go from any node to any node
- Maximum depth is $\log_2 n$ where n is the number of nodes
 - Because we require completeness and we have 2 possible children

Binary Heap

Using Complete Binary Tree to make priority_queue


Binary Heap

- We use **Complete Binary Tree** to store data
 - A **value** is stored at the **node**
- When data is modified (via push or pop), we must maintain these rules
 1. Tree must always be **Complete Binary Tree**
 2. For any node, its **value must be more than that of its children**




Getting Maximum Data

- Root contains highest value (because of rules 2.)
- Top() just simply return root




Adding data to Binary Heap

- Maintain Binary Heapness
 - structure of data
 - Value of data
- Structure rules says where the new node should be
 - Next to **right-most child** of the **deepest level**
 - But if we put the new data there, the **value rules might be broken**
 - Fix it



push(60)

Fix from adding a new node



Value rules: parent more than children


See that, after each swap, the **blue swapped node** does not violate value rules with **its new children**

Fix:


- Check where we just add a node, if value rules is broken, swap with parent
- After swap, re-check with new parent
- Keep doing until correct or at root

Delete maximum data

- Similar to push, we will try to maintain structure first
- Delete will remove root, find something to replace
 - Use the lowest, right-most node
- Value rules might be broken
 - Fix it



Fix from deleting root node



Fix:

- Start at replaced root, if value rules is broken, swap with maximum child
- After swap, re-check with new children
 - Beware! There is a case where we might have only one child
- Keep doing until correct or has no child

Value rules: parent more than children

See that, after each swap, the blue swapped node does not violate value rules with its parent and children

Analysis

- How fast is each push, pop
- Push
 - Add to a vector is $O(1)$ amortized
 - Fixing value rules is $O(h)$ where h is the maximum number of depth of the tree (we call this value tree height)
 - Notice that tree height is $O(\lg n)$
- Pop
 - Fixing value rules is $O(h)$ where h is the maximum number of depth of the tree (we call this value tree height)


$O(\log n)$

$O(\log n)$

CP::priority_queue

How to store a tree?

- Use dynamic array
 - Each node can be labelled from 0 to n-1
- Root is at 0
- Left child of node i is at $(i*2)+1$
- Right child of node i is at $(i*2)+2$
- Parent of node i is at $(i-1)/2$



0	1	2	3	4	5	6	7	8	9
75	50	13	40	3	1	10	35	36	2

Layout

```
template <typename T,typename Comp = std::less<T> >
class priority_queue {
protected:
    T *mData;
    size_t mCap;    Same as CP::vector
    size_t mSize;
    Comp mLess;
    void expand(size_t capacity) {}
    void fixUp(size_t idx) {}
    void fixDown(size_t idx) {}
public:
    //----- constructor -----
    priority_queue(priority_queue<T,Comp>& a);
    priority_queue(const Comp& c = Comp() );
    priority_queue<T,Comp>& operator=(priority_queue<T,Comp> other);
    ~priority_queue();
    //----- capacity function -----
    bool empty() const;
    size_t size() const;
    //----- access -----
    const T& top();
    //----- modifier -----
    void push(const T& element);
    void pop();
};
```

Will talk about later

Fix value rules

Constructor

```
priority_queue(const Comp& c = Comp() ) :  
    mData( new T[1]() ),  
    mCap( 1 ),  
    mSize( 0 ),  
    mLess( c )  
{ }
```

```
priority_queue(priority_queue<T,Comp>& a) :  
    mData(new T[a.mCap]()),  
    mCap(a.mCap),  
    mSize(a.mSize),  
    mLess(a.mLess)  
{  
    for (size_t i = 0; i < a.mCap;i++)  
        mData[i] = a.mData[i];  
}
```

- Using list initialize
- See that `mData` is dynamic array in the same way as vector
- `mLess` is something that is just either copied or default initialize
 - Will talk about it later

Destructor and Copy Assignment Operator

```
~priority_queue() {  
    delete [] mData;  
}
```

- Using standard copy-and-swap idiom

```
priority_queue<T,Comp>& operator=(priority_queue<T,Comp> other) {  
    using std::swap;  
    swap(mSize,other.mSize);  
    swap(mCap,other.mCap);  
    swap(mData,other.mData);  
    swap(mLess,other.mLess);  
    return *this;  
}
```

Push


```
void expand(size_t capacity) {
    T *arr = new T[capacity]();
    for (size_t i = 0; i < mSize; i++) {
        arr[i] = mData[i];
    }
    delete [] mData;
    mData = arr;
    mCap = capacity;
}
```

Same as
CP::vector

```
void push(const T& element) {
    if (mSize + 1 > mCap)
        expand(mCap * 2);
    mData[mSize] = element;
    mSize++;
    fixUp(mSize-1);
}
```

- See that the right-most child of the deepest level is at `mData[mSize-1]` and the new node should be at `mData[mSize]`
- We do the same thing as vector's `push_back`
- Then fix the value rule

Fix Up




- Instead of actual swap, we perform **insert** and **find** appropriate position at the same time

```
void fixUp(size_t idx) {
    T tmp = mData[idx];
    while (idx > 0) {
        size_t p = (idx - 1) / 2;
        if ( tmp < mData[p] ) break;
        mData[idx] = mData[p];
        idx = p;
    }
    mData[idx] = tmp;
}
```



```
void fixUp(size_t idx) {
    while (idx > 0) {
        size_t p = (idx - 1) / 2;
        if ( mData[idx] < mData[p] ) break;
        T tmp = mData[p];
        mData[p] = mData[idx];
        mData[idx] = tmp;
        idx = p;
    }
}
```

Fix Up



```
void fixUp(size_t idx) {  
    T tmp = mData[idx];  
    while (idx > 0) {  
        size_t p = (idx - 1) / 2;  
        if ( tmp < mData[p] ) break;  
        mData[idx] = mData[p];  
        idx = p;  
    }  
    mData[idx] = tmp;  
}
```

mData

0	1	2	3	4	5	6	7	8	9
75	60	13	40	50	1	10	35	36	2

mLess

- priority_queue allows a custom comparator
- Custom comparator X
 - We must be able to $X(a,b)$ where X will compare a and b and return true only when a is less than b
 - X is a **variable** that implement `operator()`

```
int main() {  
    less<int> x;  
    greater<int> y;  
  
    int a = 10;  
    int b = 3;  
    cout << x(a,b) << endl;  
    cout << y(a,b) << endl;  
}
```

mLess

- Initialize at constructor as variable `mLess` to be of type `Comp` in template
- Any comparison of our data (type `T`) must be done by `mLess`

```
template <typename T,typename Comp = std::less<T> >
class priority_queue {
    //...
    T *mData;
    size_t mCap;
    size_t mSize;
    Comp mLess;
    //...
    priority_queue(const Comp& c = Comp() ) :
        mData( new T[1]() ),
        mCap( 1 ),
        mSize( 0 ),
        mLess( c )
    { }
```

```
void fixUp(size_t idx) {
    T tmp = mData[idx];
    while (idx > 0) {
        size_t p = (idx - 1) / 2;
        if ( mLess(tmp,mData[p]) ) break;
        mData[idx] = mData[p];
        idx = p;
    }
    mData[idx] = tmp;
}
```

Pop

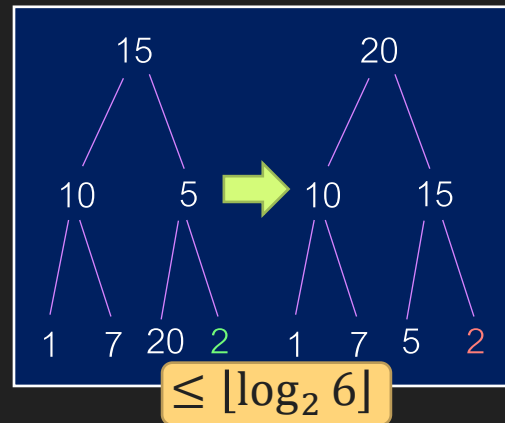
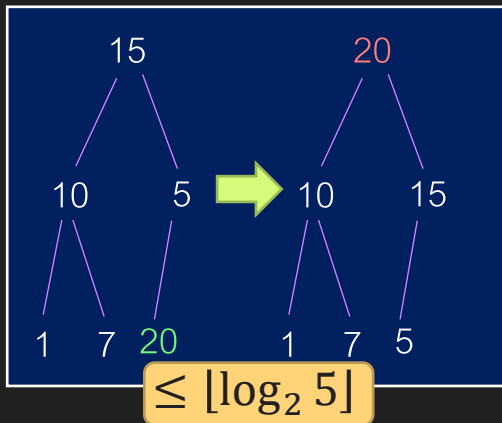
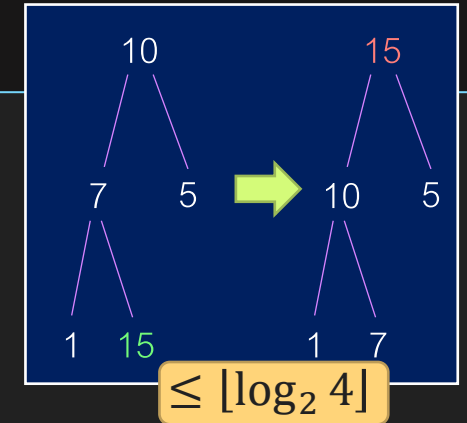
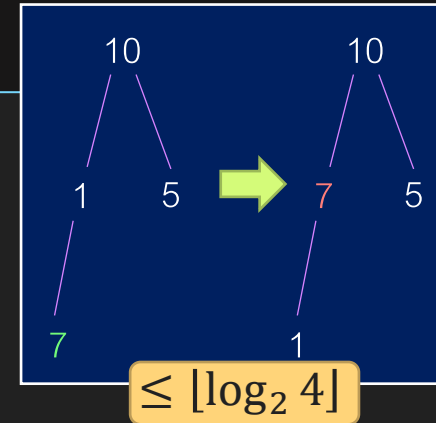



```
void pop() {  
    mData[0] = mData[mSize-1];  
    mSize--;  
    fixDown(0);  
}
```

```
void fixDown(size_t idx) {  
    T tmp = mData[idx];  
    size_t c;  
    while ((c = 2 * idx + 1) < mSize) {  
        if (c + 1 < mSize && mLess(mData[c],mData[c + 1]) ) c++;  
        if ( mLess(mData[c],tmp) ) break;  
        mData[idx] = mData[c];  
        idx = c;  
    }  
    mData[idx] = tmp;  
}
```

- While loop check if we have at least one child
- `c` is the index of highest value child
 - Must consider the case where we have only one child
- Exercise: read the rest yourself

Construct PQ from n data

```
priority_queue(std::vector<T> &v, const Comp& c = Comp() ) :
    mData( new T[v.size()]() ), mCap( v.size() ), mSize( 0 ), mLess( c ) {
    for (size_t i = 0; i < mSize; i++) push(v[i]);
}
```




$$\begin{aligned} \text{Total} &\leq \lfloor \log_2 1 \rfloor + \lfloor \log_2 1 \rfloor + \dots + \lfloor \log_2 n \rfloor \\ &\leq \lfloor \log_2 (1 \times 2 \times 3 \times \dots \times n) \rfloor = \lfloor \log_2 n! \rfloor \end{aligned}$$

$\lfloor \log_2 n! \rfloor$ is $O(n \log n)$

Better Method

```
priority_queue(std::vector<T> &v, const Comp& c = Comp() ) :  
    mData( new T[v.size()]() ), mCap( v.size() ), mSize( v.size() ), mLess( c )  
{  
    for (size_t i = 0; i < mSize; i++) mData[i] = v[i];  
    for (int i = mSize/2-1; i >= 0; i--) fixDown(i);  
}
```




- Consider each node to be a Binary Heap
- Fix down from back to front

How fast?

Total node = 31

Tree height = $\log_2 31 = 4$



Depth	nodes	Max fix per node
0	1	4
1	2	3
2	4	2
3	8	1
4	16	0

Binary Tree Property:

- There are at most 2^k nodes at depth k
- For a tree of height h , at depth k , fix down need at most $h-k$ iterations

$$= 2^h \sum_{k=0}^h k 2^{-k} < 2^h \sum_{k=0}^{\infty} k 2^{-k} = 2^h 2 = 2^{h+1} = 2^{\log_2 n + 1} = O(n)$$

This is 2

สูง จำนวนต้น

4 × 1

3 × 2

2 × 4

1 × 8

0 × 16

$k \times 2^{h-k}$

h

$\leq \sum_{k=0}^h k 2^{h-k}$